

L1

Q. What is Data -

- data is a collection of raw, unorganized facts and details like text, observations, figures, symbols, and descriptions of things.
- in other words, data does not carry any specific purpose and has no significance by itself

- moreover, data is measured in terms of bits and bytes - which are basic units of information in the context of computer storage and processing
- data can be recorded and doesn't have any meaning unless processed

Q. Types of data

- o Quantitative
 - numerical form
 - weight, volume, cost of an item

- o Qualitative
 - descriptive, but not numerical
 - name, gender, hair color of a person

Q. What is Information

- information is processed, organized, and structured data
- it provides context of the data and enables decision making
- processed data that make sense to us
- information is extracted from the data, by analyzing and interpreting piece of data

e.g. - you have data of all the people living in your locality. Its data when you analyse and interpret the data and come to some conclusion that

- there are 1000000 citizens
- the sex ratio is 1.1
- these are information
- new born babies are 100

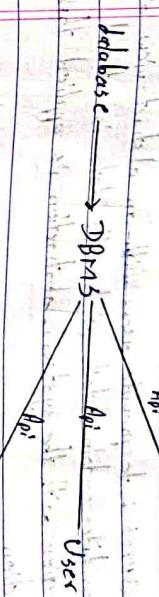
□ Data vs information

- data is collection of facts, while information puts those facts into context
- while data is raw and unorganized, information is organized.
- data points are individual and sometimes unrelated, information maps out what data to provide a big-picture view of how it all fits together
- data, on its own, is meaningless. When it's analyzed and interpreted, it becomes meaningful information.
- data does not depend on information. However, information depends on data
- data typically comes in the form of graphs, numbers, figures, or statistics. Information is typically presented through words, language, thoughts, and ideas.
- data isn't sufficient for decision-making but you can make decisions based on information.

□ DBMS vs File system

- File-processing system has major disadvantages
 - o data redundancy and inconsistency
 - o difficulty in accessing data
 - o data isolation
- Integrity problems
- Atomicity problems
- Concurrent access anomalies
- Security problems

□ Database



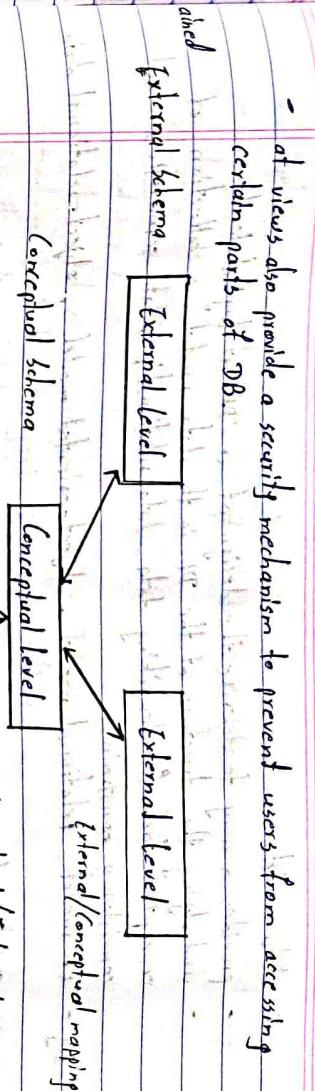
□ DBMS

- o What is database
- database is an electronic place/system where data is stored in a way that it can be easily accessed, managed, and updated
- To make real use of data, we need database (DBMS)
- o What is DBMS
- a database management system is a collection of interrelated data and a set of programs to access those data. The collection of data, usually referred to as the database, contains information relevant to an enterprise. The primary goal of a DBMS is to provide a way to store and retrieve database information that is both convenient and efficient.
- A DBMS is the database itself, along with all the software and functionality it is used to perform different operations, like addition, access, updating and deletion of the data.

DBMS Architecture

a) View of data (Three schema architecture)

- the major purpose of DBMS is to provide users with an abstract view of the data, that is, the system hides certain details of how the data is stored and maintained.
- to simplify user interaction with the system, abstraction is applied through several levels of abstraction.
- the main objective of three level architecture is to enable multiple users to access the same data with a personalized view while skipping the underlying data only once.



Conceptual Schema

Conceptual Level

External/Internal mapping

Internal Schema

Internal Level

DB

Conceptual / Internal mapping

Conceptual Level

External/Internal mapping

External Schema

External Level

External Schema</

- provides a way to describe the design of a DB at logical level
- underlying the structure of the DB is the data model, a collection of conceptual tools for describing data, data relationship, data semantics and consistency constraints

Eg - ER model, relational model, object-oriented model, object-relational data

ii Database languages

- Data Definition Language (DDL) to specify the database schema.
- Data Manipulation language (DML) to express database queries and updates
- practically : both language features are present in a single DB language.

Eg - SQL language

iii DDL

- we specify consistency constraints, which must be checked, every time DB is updated.

iv DML

- Data manipulation involves

- retrieval of information stored in DB
- insertion of new information into DB
- deletion of information from the DB
- updating existing information stored in DB

Query language - a part of DML to specify statement requesting the retrieval of information

Eg - Database accessed from Application program's perspective

Eg - Apps (written in host languages, C/C++, Java...) interacts with DB

- Banking system's module generating payroll access DB by executing DML statements from the host language
- API is provided to send DML/DDL statements to DB and retrieve the results
- o Open database connectivity (ODBC), microsoft, "C"
- o Java database connectivity (JDBC), Java

Database Administrator, DBA

- a person who has central control of both the data the programs that access those data

o Functions of DBA

- schema definition
- storage structure and access methods
- schema and physical organization modifications
- authorization control
- routine maintenance
 - periodic backups
 - security patches
 - any upgrades

ii DBMS application architectures

- client machines, on which remote DB users work, and server machines on which DB system runs

① T1 architecture

- the client, server and DB all present on the same machine

② T2 architecture

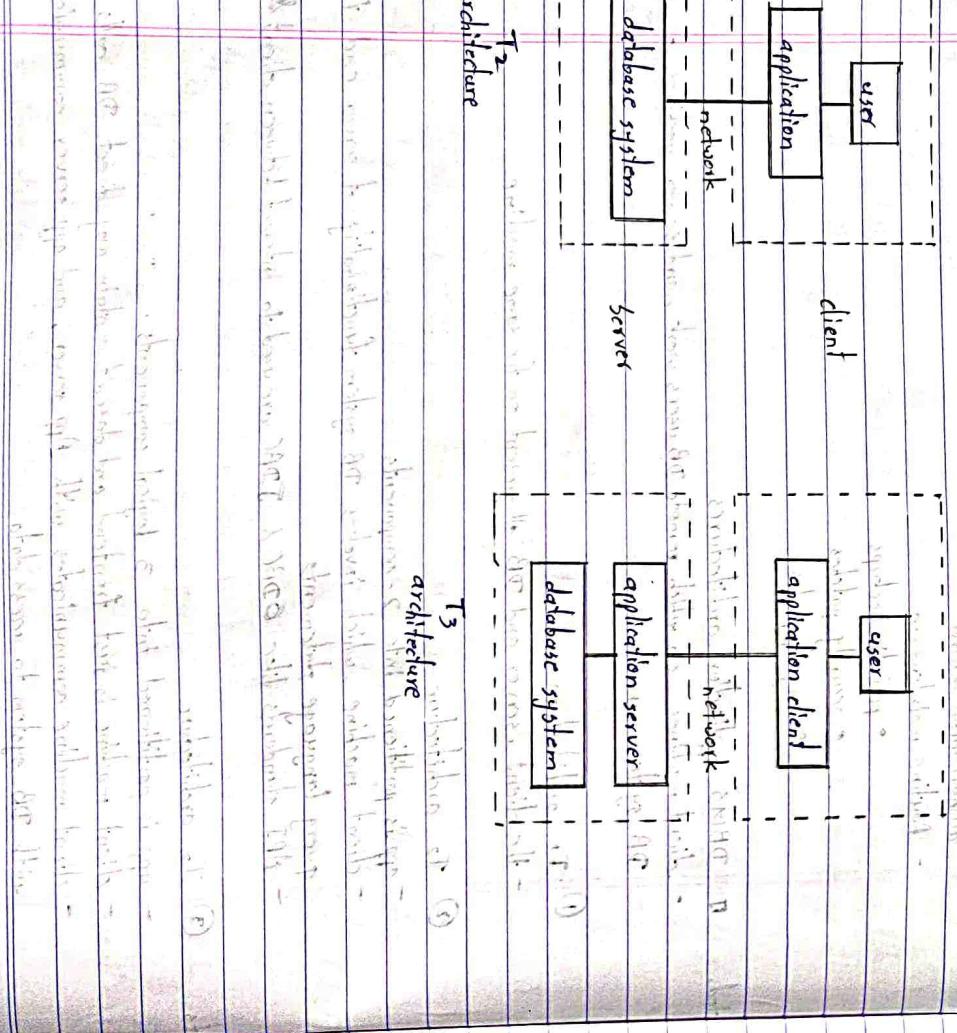
- app is partitioned into 2-components
- client machine, which invokes DB system functionality at server end through query language statements
- API standards like ODBC & JDBC are used to interact between client & server

③ T3 architecture

- app is partitioned into 3 logical components
- client machine is just frontend and doesn't contain any direct DB calls
- client machine communicates with App server, and app server communicated with DB system to access data.

Entity-Relationship Model

- business logic, what action to take at that condition is in app server itself.
- T₃ architecture are best for web applications
- Advantages
 - scalability due to distributed application server
 - data integrity - app server acts as a middle layer between client and DB which minimize the chances of data corruption
 - security - client can't directly access DB, hence it is more secure



13

- data model - collection of conceptual tools for describing data, data relationship, data semantics, and consistency constraints.
- ER model -
 - it is high level data model based on a perception of a real world that consists of a collection of basic objects, called entities, and of relationships among these objects
 - graphical representation of ER model is ER diagram, which acts as a blueprint of DB

Entity -

- an entity is a "thing" or "object" in the real world that is distinguishable from other all objects.
- it has physical existence
- each student in a collage is an entity
- entity can be uniquely identified (by a primary attribute, aka primary key)
- **Strong entity** - can be uniquely identified
- **Weak entity** - can't be uniquely identified, depends on some other strong entity
 - it doesn't have sufficient attributes, so select a uniquely identifiable attribute

T₂

architecture

T₃

architecture

- sequential number counter can be generated separate for each loan
- weak entity depends on strong entity for existence
- Entity set
 - it is a set of entities of the same type that share the same properties, or attributes.

Eg - student is an entity set

Eg - customer of a bank

Notes on Entity Relationship

⑤

Multi-valued

- attribute having more than one value
- phone number, nomine-name on some insurance, dependent name etc.
- limit constraint may be applied, upper or lower limits

Eg -

- student entity has following attributes.

- o name

- o standard

- o course

- o batch

- o contact number

- o address

Types of Attributes

⑥ Simple

- attributes which can't divide further

Eg - customer's account number in a bank, student's roll number, etc.

Composite

- can be divided into subparts (that is other attributes)

Eg - name of person, can be divided into first-name, middle-name, last-name

- if used wants to refer to an entire attribute or to only a component of the attribute

- address can also be divided, street, city, state, pin code

Single-value

- only one value attribute

Eg - student ID, loan number for loan

Relationship

- association among two or more entity

Eg - person has vehicle, parent has child, customer borrows loan, etc.

- strong relationship, between two independent entities
- weak relationship between weak entity and its owner/strong entity

Eg - loan instalment - payments > payment

degree of relationship

- number of entities participating in relationship

Unary - only one entity participates. Eg - employee manages employee

Binary - two entities participates. Eg - student take course

Ternary - three entities participates. Eg - employee works-on branch, employee works-on job

- binary and common

a Relationship constraints

- mapping cardinality / cardinality ratio
- number of entities to which another entity can be associated via a relationship.

* one-to-one entity in A associates with at most one entity in B, where A & B are entity sets, and an entity of B is associated with at most one entity of A

Eg - citizen has address card

* one-to-many entity in A associates with at most one entity in B, while entity in B is associated with at most one entity in A

Eg - citizen has vehicle

* Many-to-one entity in A associates with at most one entity in B, while entity in B can be associated with N entity in A

Eg - course taken by professor

* many-to-many entity in A associates with N entity in B, while entity in B also associates with N entity in A

Eg - customer buy product

- student attend course

b Participation constraints

1. Alpha, minimum cardinality constraint

2. Beta, maximum cardinality constraint

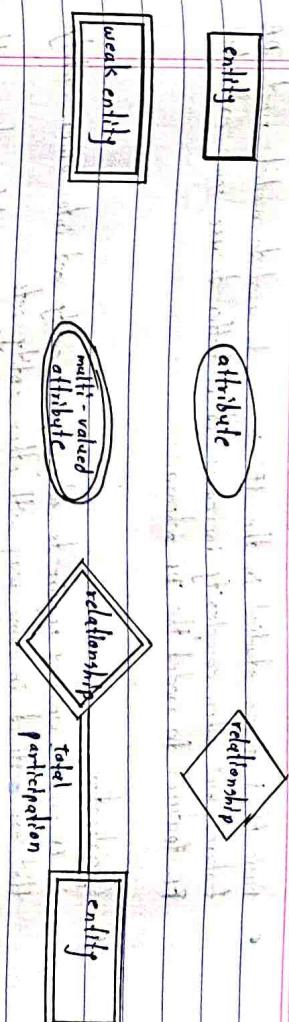
- partial participation, not all entities are involved in the relationship instance

- total participation, each entity must be involved in at least one relationship instance

Eg - customer borrow loan, loan has total participation as it can't exist without customer entity and customer has partial participation

- weak entity has total participation constraint, but strong may not have total

c ER notations



Extended ER Features

- Basic ER features, studied in the EER-3, can be used to model most DB features but when complexity increases, it is better to use some extended ER features to model the DB schema.

a) Specialisation

- in ER model, we may require to subgroup an entity set into other entity sets that are distinct in some way with other entity sets
- specialisation is splitting up the entity set into further sub entity sets on the basis of their functionalities, specialities and features

Eg - it is a Top-Down approach

- person entity set can be divided into customer, student, employee, person is superclass and other specialised entity sets are subclasses
- we have "is-a" relationship between superclass and subclass

depicted by triangle component

b) Why Specialisation

- both generalisation and specialisation, has attribute inheritance.
- the attribute of higher level entity sets are inherited by lower level entity sets.

Eg - customer and employee inherit the attributes of person.

c) Participation Inheritance

- if a parent entity set participates in relationship then its child entity sets will also participate in that relationship.

d) Aggregation

- how to show relationships among relationships? aggregation is the technique
- abstraction is applied to treat relationships as higher-level entities. we can call it abstract entity
- avoid redundancy by aggregating relationship as an entity set itself

e) Generalisation

- it is just a reverse of specialisation
- DB designer may encounter certain properties of two entities are overlapping, designer may consider to make a new generalised entity set. that generalised entity set will be a super class
- "is-a" relationship is present between subclass and super class
- car, jeep and bus all have some common attributes, to avoid data representation for the common attributes, DB designer may consider to generalise to a new entity set "vehicle"
- it is a Bottom-up approach

f) Why Generalisation

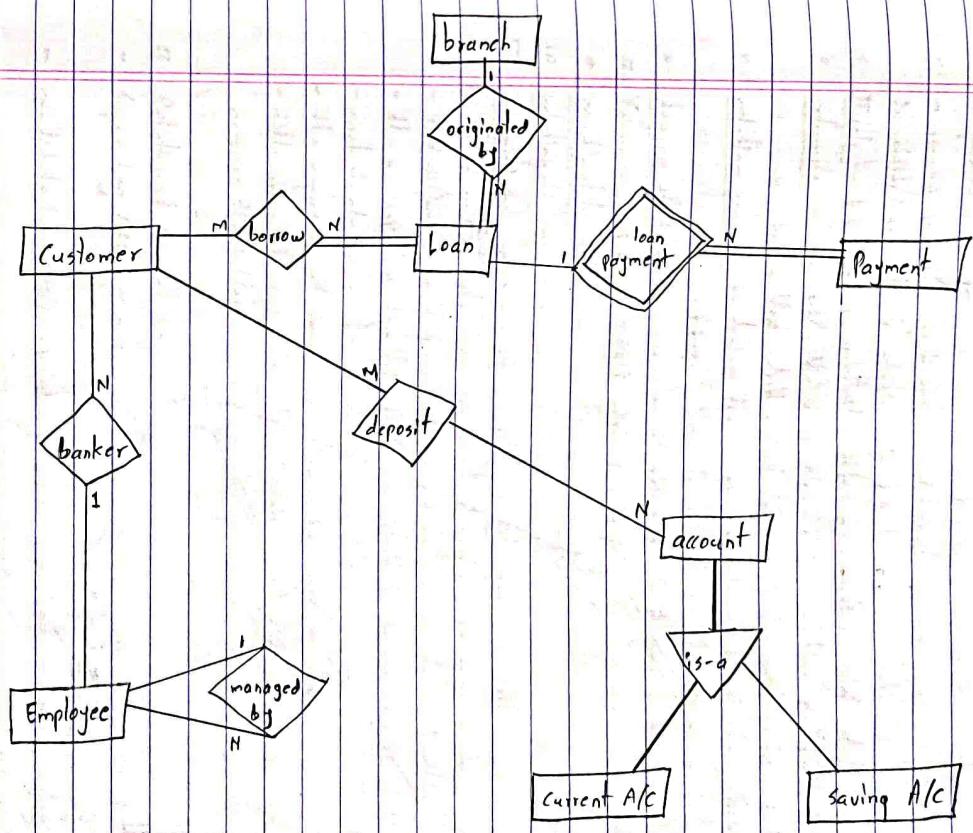
- make DB more refined and simpler
- common attributes are not repeated

steps to make ER diagram

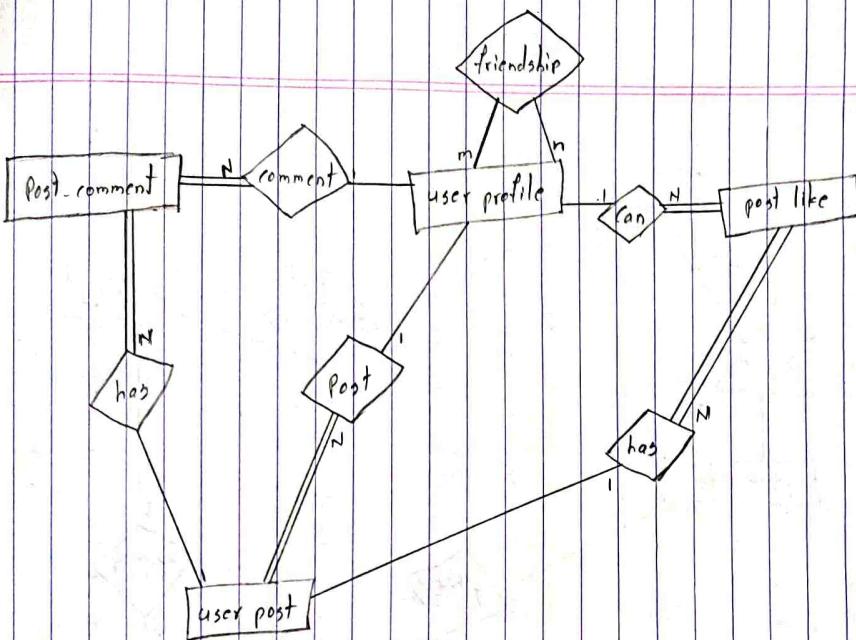
- ① Identify Entity sets
- ② identifying attributes and their types
- ③ Relationship and constraints

→ mapping
participation

every entity have their attributes



every entity have attribute



L7

Relational model

- relational model (RM) organizes the data in the form of relations (Tables)
- a relational DB consists of collection of tables, each of which is assigned a unique name
- a row in a table represents a relationship among a set of values, and table is collection of such relationships
- Tuple - a single row of the table representing single data point / a unique record
- columns - represents the attributes of the relation, each attribute, there is a permitted value, called domain of the attribute
- relation schema - defines the design and structure of the relation, contains the name of the relation and all the columns/attributes
- degree of table - number of attributes/columns in a given table/relation
- common RM based DBMS systems, aka RDBMS : oracle, IBM, MySQL, MS Access
- cardinality - total no. of tuples in a given relation
- candidate key - set of attributes which can uniquely identify an each tuple
- impostent properties of a table in relational model
 - the name of relation is distinct among all other relation
 - the values have to be atomic, can't be broken down further
 - the name of each attribute/column must be unique
 - each tuple must be unique in a table
 - the sequence of row and column has no significance
 - table must follow integrity constraints - it helps to maintain data consistency across the tables
- **Relational Model keys**
 - super key (SK) - any PK or attributes present in a table which can uniquely identify each tuple
 - candidate key (CK) - minimum subset of superkeys, which can uniquely identify each tuple, it contains no redundant attribute
 - CK value shouldn't be null
 - primary key (PK) - selected out of CK set, has the least no. of attribute

- all candidate key (AK) - all CK except PK
- foreign key (FK) -
 - it creates relation between two tables
 - a relation, say r_1 may include among its attributes the PK of an other relation, say r_2 . this attribute is called FK from r_1 referencing r_2 .
 - the relation r_1 is aka referencing (child) relation of the RR dependency, and r_2 is called referenced (parent) relation of the FK.
 - FK helps to cross reference between two different relations
 - Composite Key - PK formed using at least 2 attributes
 - Compound Key - PK which is formed using 2 FK
 - Surrogate Key - synthetic PK generated automatically by DB, usually an integer value
 - may be used as PK
- Integrity constraints
 - CRUD operations must be done with some integrity policy so that DB is always constraint introduced so that we do not accidentally corrupt the DB
 - Domain constraints
 - restricts the value in the attribute of relation, specifies the domain
 - restrict the data types of every attribute
 - e.g. - we want to specify that the enrolment should happen for candidate birth year < 2002.
 - Key constraints - the six types of key constraints present in the database management system are -
 - NOT NULL - this constraint will restrict the user from not having Null value. it ensures that every element in the database has a value
 - UNIQUE - it helps us to ensure that all the values consisting in a column are different from each other
 - DEFAULT - it is used to set the default value to the column. the default value is added to the columns. if no value is specified for them
 - CHECK - it is one of the integrity constraints in DBMS. it keeps the check that integrity of data is maintained before the end after the completion of the CRUD
 - PRIMARY KEY - this is an attribute or set of attribute that can uniquely identify each entity in the entity set. the primary key must contain unique as well as not null values
 - FOREIGN KEY - whenever there is relationship between two entities, there must be some common attributes between them. this common attributes must be the primary key of one entity set and will become the foreign key of another entity set. this key will prevent every action which can result in loss of connection between tables
- Entity constraints
 - every relation should have PK, PK != NULL
- Referential constraints
 - specified between two relations & helps maintains consistency among tuples of two relations

Transform ER model to relational model

- Both ER model and relational model are abstract logical representation of real-world enterprises. because the two models implies the similar design principles, we can convert ER model design into relational design.
- Considering a DB representation from an ER diagram has a table format is the way we arrive at relational DB design from an ER diagram
- ER diagram notations to relations

Strong Entities

- becomes an individual table with entity name, attributes, becomes columns of the relation
- entities primary key (PK) is used as relation's PK
- FK are added to establish relationships with other relations

Weak Entity

- a table is formed with all the attributes of an entity
- PK of its corresponding strong entity will be added as FK
- PK of the relation will be a composite PK {FK + Partial Discriminator Key}

- Q Single values Attributes
- represented as columns directly in the tables/relations

Composite Attributes

- handled by creating a separate attribute itself in the original relation for each composite attribute

- Eg - Address {street-name, house-no} is a composite attribute in customer relation, we add address-street-name & address-house-name as new columns in the attribute and ignore "address" as an attribute

Multivalued Attributes

- new tables (named as original attribute name) are created for each multivalued attribute



- PK of the entity is used as column FK in the new table
- multivalued attribute's similar name is added as column to define multiple values
- PK of the new table would be {FK + multivalued-name?}
- Eg - For strong entity Employee, dependent-name is a multivalued attribute
 - o new table named dependent-name will be formed with columns emp-id & name
 - o PK {emp-id, name?}
 - o FK {emp-id}

Q) Derived Attributes - not considered in the tables

Q) Generalisation

- Method 1
 - create a table for the higher level entity set. for each lower-level entity sets, create a table that includes a column for each of the attributes of that entity set plus a column for each attribute of the primary key of the higher-level entity set
 - Table 1 - account (account-number, balance)
- Table 2 - saving-account (account-number, interest-rate, daily-withdrawl-limit)
- Table 3 - current-account (account-number, overdraft-amount, per-transaction-charges)

Method 2

- an alternative representation is possible, if the generalisation is disjoint and complete - that is, if no entity is a member of two lower-level entity sets directly below a higher-level entity set, and if every entity in the higher level entity set is also a member of one of the lower-level entity sets, here do not create a table for the higher-level entity set instead, for each lower-level entity set, create a table that includes a column for each of the attributes of that entity set plus a column for each attribute of the higher-level entity sets
- Table would be :

- Table 1 - saving account (account-number, balance, interest-rate, daily-withdrawl-limit)
- Table 2 - current account (account-number, balance, overdraft-amount, per-transaction-charges)

Q) Aggregation

- table of the relationship set is made
- attributes includes primary key of entity set and aggregation set's entities
- also add descriptive attribute of it any on the relationship.

for example

Drawbacks of method 2

- if the second method were used for an overlapping generalisation, some values such as balance could be stored twice unnecessary similarly, if the generalisation, some values such as some accounts were neither saving nor current accounts - then such accounts could not be represented with the second method.

SQL

DATA TYPE	Description
CHAR	string (e-255) . string with size = (0-255) eg - CHAR(25)
VARCHAR	string (0-255)
TINYTEXT	string (0-255)
TEXT	string (0 - 65535)
LARGE BLOB	string (0 - 65535)
MEDIUM BLOB	string (0 - 16777215)
LONGTEXT	string (0 - 4294967295)
CONGBLOB	string (0 - 4294967295)
TINYINT	integer (-128 to 127)
SMALLINT	integer (-32768 to 32767)
MEDIUMINT	integer (-8388608 to 8388607)
INT	integer (-2147483648 to 2147483647)
FLOAT	decimal with precision 10-23 digits
DOUBLE	decimal with 12-15 digits

DECIMAL	double stored as string
DATE	YYYY - MM - DD
DATETIME	YYYYMMDD HHMMSS
TIMESTAMP	HH:MM:SS
TIME	one of the preset values
ENUM	one or many of the preset values
SET	0/1
BOOLEAN	cg - BIT(n), n upto 64, store values in bits
BIT	

- size : TINY < SMALL < MEDIUM < INT < BIGINT
- variable length data types: Eg - VARCHAR are better to use because it uses only memory which ever are needed

- values can also be unsigned. Eg - INT UNSIGNED

Types of SQL commands

- DDL data define language (define relation schema)
 - CREATE - create Table, DB, view
 - ALTER TABLE - modification in table structure Eg - change column datatype / add, remove
 - DROP - delete Table, DB, view
 - TRUNCATE - remove all the tuples from the table
 - RENAME - rename DB name, table name, column name etc.
- DML / DQL data retrieval language / data query language (retrieve data from tables) Eg -
 - SELECT
- DML data modification language (use to perform modification in DB)
 - INSERT - insert data into a relation
 - UPDATE - update relation data
 - DELETE - delete row(s) from the relation

4. DCL data control language (grant or revoke authorities from user)

- GRANT - access privileges to the DB
- REVOKE - revoke user access privileges

TCL

translation control language (to manage transactions done in the DB)

- START TRANSACTION - begin a transaction
- COMMIT - apply all the changes and end transaction
- ROLLBACK - discard changes and end transaction
- SAVEPOINT - checkout within the group of transaction in which to rollback

MANAGING DB (DDL)

- creation of DB - CREATE DATABASE IF NOT EXISTS db-name;
- USE db-name; To choose DB
- DROP DATABASE IF EXISTS db-name; To delete database
- SHOW DATABASES; list all the DBs in the server
- SHOW TABLES; list tables in the selected DB

DATA RETRIEVAL LANGUAGE (DQL)

- syntax - SELECT <set of column name> From <table names>;
- order of execution RIGHT TO LEFT
- can use SELECT keyword without using From clause?
- Yes, using DUAL Tables
- DUAL tables are dummy tables created by MySQL, help users to do certain obvious actions without referring to user defined Tables

- SELECT 55+11;
- SELECT now();
- SELECT user(); / session();

o WHERE

- reduce rows based on given conditions.

Eg -

- `SELECT * FROM customer WHERE age > 18;`

o BETWEEN

- Eg - `SELECT * FROM customer WHERE age between 20 AND 100;`

o IN

- reduces OR conditions

- Eg - `SELECT * FROM officers WHERE officer_name IN ('Com', 'RAM', 'xx');`

o AND /OR /NOT

- AND - WHERE cond1 AND cond2
- OR - WHERE cond1 OR cond2
- NOT - WHERE column NOT IN (1,2,3,4);

o IS NULL

- Eg - `SELECT * FROM customer WHERE prime_Status IS NULL;`

o Pattern Searching / Wildcard ('%', '_')

- '%' any number of character from 0 to n similar to '*' asterisk in regex.
- '_' any one character

o ORDER BY

- sorting the data retrieved using WHERE clause.
- `ORDER BY <column-name> DESC;` / ASC for ascending
- Eg - `SELECT * FROM customer ORDER BY name DESC;`

o GROUP BY

- Group By clause is used to collect data from multiple records and group the result by one or more column(s). It is generally used in SELECT statement.
- Groups into category based on column given.
- `SELECT c1, c2, c3 FROM sample_table WHERE cond GROUP BY c1, c2, c3;`
- all the column names mentioned after SELECT statement shall be replaced in GROUP BY. In order to successfully execute the query.
- used with aggregation functions to perform various actions.

o COUNT()

o SUM()

o AVG()

o MIN()

o DISTINCT

- find distinct values in the table

- `SELECT DISTINCT column_name FROM table-name;`

- GROUP BY can also be used for the same

- `SELECT column_name FROM table-name GROUP BY column_name;` same diff as above

- SQL is smart enough to realize that if you are using GROUP BY and not using any aggregation function, then you mean "DISTINCT".

o GROUP BY HAVING

- out of the categories made by GROUP BY, we would like to know only particular thing ('cond')

o similar to WHERE

- `SELECT COUNT(cust_id), country FROM customer GROUP BY country HAVING COUNT(cust_id) > 50;`

* WHERE vs HAVING

- both have same function of filtering the row base on certain conditions

- WHERE clause is used to filter the rows from the table based on specified condition

3) HAVING clause is used to filter the rows from the groups based on the specified conditions

4) HAVING is used after GROUP BY while WHERE is used before GROUP BY clause

5) if you are using HAVING, GROUP BY is necessary

6) WHERE can be used with SELECT, UPDATE & DELETE keywords while GROUP BY is used with SELECT.

```
CREATE DATABASE ORG;
USE ORG;
```

```
SHOW DATABASES;
```

```
USE ORG;
```

```
CREATE TABLE worker(
```

```
WORKER_ID INT NOT NULL PRIMARY KEY AUTO_INCREMENT,
```

```
FIRST_NAME VARCHAR(255),
```

```
LAST_NAME VARCHAR(255),
```

```
SALARY INT(5),
```

```
JOINING_DATE DATETIME,
```

```
DEPARTMENT VARCHAR(255)
```

```
);
```

```
INSERT INTO worker VALUES ('1', 'om', 'KALDATE', 10000, '11-12-13 11.12.01', 'HR');
```

```
INSERT INTO worker(worker_id, first_name, last_name, salary, joining_date, department)
```

```
(2, 'RAM', 'KALDATE', 100, '11-12-13 10.12.10', 'CEO');
```

```
(3, 'GANDHI', 'MAHAN', 150, '11-02-13 10.12.10', 'CEO');
```

```
(4, 'HANNA', 'SHAWAR', 2000, '12-02-13 11.12.10', 'CEO');
```

```
(5, 'ABU', 'KALDATE', 2500, '13-02-13 11.13.11', 'MANAGER');
```

```
SELECT * FROM worker
```

```
SELECT * FROM worker WHERE FIRST_NAME LIKE 'M%';
```

```
SELECT * FROM worker WHERE FIRST_NAME LIKE '_M%';
```

```
SELECT * FROM worker WHERE FIRST_NAME LIKE '_M_';
```

SELECT * FROM worker // show entire table

WORKER_ID	FIRST_NAME	LAST_NAME	SALARY	JOINING_DATE	DEPARTMENT
1	om	KALDATE	10000	11-12-13 11.12.01	HR
2	RAM	KALDATE	100	11-12-13 10.12.10	CEO
3	GANDHI	MAHAN	150	11-02-13 11.12.10	CEO
4	HANNA	SHAWAR	2000	12-02-13 11.13.11	MANAGER
5	ABU	KALDATE	2500	13-02-13 11.13.11	EMPLOYEE

SELECT WORKER_ID, DEPARTMENT FROM worker; shows only column of WORKER_ID & DEPARTMENT

SELECT * FROM worker WHERE salary < 200; shows all rows satisfy which

SELECT * FROM worker WHERE DEPARTMENT = 'HR'; shows 1 ID's row

SELECT * FROM worker WHERE salary BETWEEN 100 AND 200; shows 3 ID's

SELECT * FROM worker WHERE salary = 100 OR salary = 150;

SELECT * FROM worker WHERE DEPARTMENT IN ('HR', 'CEO');

all the HR's & CEO's are printed

SELECT * FROM worker WHERE salary IS NULL; null salary employee prints

SELECT * FROM worker WHERE FIRST_NAME LIKE '_M'; # LIKE is IMP

→ only one character % → any number of character from 0 to n

→ first character must be ;

SELECT * FROM worker ORDER BY salary; → defaultly ascending (ASC)

SELECT * FROM worker ORDER BY salary DESC; descending (DESC)

SELECT DISTINCT DEPARTMENT FROM worker; shows distinct

SELECT DEPARTMENT, COUNT(DEPARTMENT) FROM worker GROUP BY DEPARTMENT;

SELECT DEPARTMENT, AVG(SALARY) FROM worker GROUP BY DEPARTMENT;

also can use MIN, MAX with Group BY

SELECT DEPARTMENT, COUNT(DEPARTMENT) FROM worker GROUP BY DEPARTMENT HAVING COUNT(DEPARTMENT) > 1;

Constraints 201

- Primary key

```

CREATE TABLE customer (
    id INT PRIMARY KEY,
    branch_id INT,
    first_name CHAR(50),
    last_name CHAR(50),
    DOB DATE,
    gender CHAR(6),
    PRIMARY KEY(id)
);

```

PK is not null, unique & only one per table.

Foreign key

FK refers to PK of other table

each relation can have any number of FK

CREATE TABLE ORDER (

 - id INT PRIMARY KEY,
 - delivery-date DATE,
 - order-placed-date DATE,
 - cust-id INT,
 - FOREIGN KEY(cust-id) REFERENCES customer(id)

);

CREATE TABLE customer (

email VARCHAR(100), unique, must be valid email address

);

UNIQUE

unique, can be null, table can have multiple unique attributes

14

- a CREATE TABLE customer (

CONSTRAINT age_check CHECK(age > 12),

- "age-check" can also avoid this, MySQL generates name of constraint automatically

2) MODIFY

- changes data type of an attribute
- ALTER TABLE table-name MODIFY col-name col-data-type;
- VARCHAR TO CHAR

ALTER TABLE customer MODIFY name CHAR(100);

3) CHANGE COLUMN

- rename column name
- ALTER TABLE table-name CHANGE COLUMN old-col new-col new-data-type;
- ALTER TABLE customer CHANGE COLUMN name my-name VARCHAR(100);

5)

DROP COLUMN

- drop a column completely
- ALTER TABLE table-name DROP COLUMN col-name;

Eg - ALTER TABLE customer DROP COLUMN middle-name;

6) RENAME

- rename table name itself.

- ALTER TABLE table-name RENAME TO new-table-name;

Eg - ALTER TABLE customer RENAME TO customer_details;

DATA MANIPULATION LANGUAGE (DML)

① Insert

- INSERT INTO table-name (col1,col2,col3) VALUES (v₁,v₂,v₃), (V₄,V₅,V₆),

② Update

- UPDATE table-name SET col₁=l₁, col₂=l₂ WHERE id = l₁

③ update multiple rows

Eg - UPDATE student SET standard = Standard + 1;

* ON UPDATE CASCADE

- can be added to the table while creating constraints. suppose there is a situation where we have two tables such that primary key of one table is the foreign key for another table if we update the primary key of the first table then using the ON UPDATE CASCADE foreign key of the second table automatically get updated

(3) Delete

- DELETE FROM table-name WHERE id=1;

- DELETE FROM table-name ;

④

DELETE CASCADE - (to overcome DELETE constraint of referential constraint)

- what would happen to child entry if parent table's entry is deleted?

CREATE TABLE ORDER(

order_id int PRIMARY KEY,
customer_id int,
delivery_date DATE;

cust_id INT,
FOREIGN KEY (cust_id) REFERENCES customer (id) ON DELETE CASCADE

CREATE TABLE ORDER(

order_id int PRIMARY KEY,
customer_id INT,
delivery_date DATE,

FOREIGN KEY (customer_id) REFERENCES customer (id) ON DELETE SET NULL

REPLACE

- primarily used for already present tuple in a table
- As UPDATE using REPLACE with the help of WHERE clause in PK, then that row will be replaced.

As INSERT, if there is no duplicate data new tuple will be inserted.

- REPLACE INTO student (id, class) VALUES (4, 3);

- REPLACE INTO table SET col₁=val₁, col₂=val₂;

```

CREATE TABLE order-details(
    order-id INT PRIMARY KEY,
    delivery-date DATE,
    cust-id INT,
    FOREIGN KEY(cust-id) REFERENCES customer(id)
);

INSERT INTO order-details VALUES
(100, '2024-01-01', 1),
(101, '2024-01-02', 1),
(102, '2024-01-02', 2);

INSERT INTO order-details VALUES (102, '2024-01-02', 2);

ALTER TABLE order-details CHANGE column delivery-date kadhe-pahaj VARCHAR(255);

ALTER TABLE order-details DROP COLUMN interest;

ALTER TABLE order-details ADD interest FLOAT NOT NULL DEFAULT 0;

ALTER TABLE order-details MODIFY interest INT NOT NULL DEFAULT 0;

DESC order-details;

ALTER TABLE order-details CHANGE column delivery-date kadhe-pahaj VARCHAR(255);

ALTER TABLE order-details DROP column interest;

ALTER TABLE order-details RENAME TO account-details;

SET SQL_SAFE_UPDATES = 0;

UPDATE order-details SET delivery-date = '2024-01-00' WHERE order-id = 100;

UPDATE customer SET pincode = 10000;

UPDATE customer SET pincode = pincode + 1;

DELETE FROM order-details WHERE order-id = 100;

All Table will delete values.

```

first customer table implement then order-id table implement

CREATE TABLE customer					
	id	c-name	gender	city	pincode
1	01	M	SAWARAJON	413624	
2	02	M	SAWARIKON	413624	
3	03	F	SOLAPUR	411004	

customer table implemented

order-id table implement

o_id	o_date	o_time	o_total	o_status
1	2023-10-10	14:30:00	10000	Shipped

order-id table implemented

* Referential Constraint

- ① INSERT
- ② DELETE constraint

↳ On delete cascade
↳ On delete Set Null

CREATE TABLE cust (

id INT PRIMARY KEY,
l-name VARCHAR(255),

gender CHAR(1),

INSERT INTO cust VALUES (1, 'RAM', 'M', 'male');
INSERT INTO cust VALUES (2, 'SHEK', 'F', 'female');

DELETE FROM cust WHERE id = 1;



CREATE TABLE ordetails (

order_id INT PRIMARY KEY,
cust_id INT,

FOREIGN KEY(cust_id) REFERENCES cust(id) ON DELETE CASCADE

);
INSERT INTO ordetails VALUES (100, 1, 1);

DELETE FROM cust WHERE id = 1;

Replace

- ↳ if data already present, replace it
- ↳ if data not present, insert it.

REPLACE INTO cust VALUES (5, 'RAM', 'M', 'male');
REPLACE INTO cust VALUES (5, 'SEX', 'F', 'female'); // id is PK

REPLACE INTO cust (id, l-name) VALUES (5, 'SAM');

REPLACE INTO cust SET id = 5, l-name = 'random';

REPLACE INTO cust(id, l-name, gender) SELECT id, l-name, gender FROM cust WHERE id = 2;

JOINING TABLES

- all RDBMS are relational in nature, we refer to other tables to get meaningful outcomes.
- FK are used to do reference to other table.

* INNER JOIN

→ returns a resultant table that has matching values from both the tables

ON DELETE CASCADE or all the tables

ON DELETE SET NULL

ON condition

INNER JOIN table_3 ON condition_2

* Aliasing MySQL (AS)

↳ Alias in MySQL is used to give a temporary name to table or column in a table for the purpose of particular query. It works as nickname, it works as a nickname for expressing the tables or column names. It makes the query short & neat.

OPEN here the details related to customer when parent is deleted the Foreign key on the child will filled by null. data remains in the table

Q) OUTER JOIN



C.J



Q) INNER JOIN

L.J

R.J

F.J

C.J

- this returns a resulting table that all the data from left-table and the matched data from right-table
- o SELECT columns FROM table1 LEFT JOIN table2 ON [join-condition];

R.J

C.J

- this returns a resulting table that all the data from right-table and the matched data from left-table
- o SELECT columns FROM table1 RIGHT JOIN table2 ON [join-condition];

F.J

C.J

- this returns a resulting table that contains all the data when there is a match on left or right table data.
- o Emulated in MySQL using LEFT and RIGHT JOIN
- o LEFT JOIN UNION RIGHT JOIN

3) FULL JOIN

C.J

- this returns a resulting table that contains all the data when there is a match on left or right table data.
- o Emulated in MySQL using LEFT and RIGHT JOIN

- o SELECT columns FROM table1 AS t1 LEFT JOIN table2 AS t2 ON t1.id=t2.id

UNION

C.J

- o SELECT columns FROM table1 AS t1 RIGHT JOIN table2 AS t2 ON t1.id=t2.id

- * UNION ALL can also be used this will duplicate values as well while UNION gives unique values

4) CROSS JOIN

C.J

- this returns all the cartesian products of the data present in both tables.
- hence all possible variations are reflected in the output

- used rarely in practical purpose

- table1 has 10 rows and table2 has 5, then resultant would have 50 rows
- o SELECT column_list FROM table1 CROSS JOIN table2;

5) SELF JOIN

C.J

- it is used to get the output from a particular table when the same table is joined to itself.
- o Emulated using INNER JOIN
- o SELECT columns FROM table1 AS t1 INNER JOIN table2 AS t2 ON t1.id=t2.id
- o John without using JOIN keywords
- o SELECT * FROM table1,table2 WHERE condition
- o SELECT artistname,albumname,year-recorded FROM artist, album WHERE artist.id = album.artist_id;
- o always gives distinct rows.
- o SET operation.
- o combines multiple tables based on matching combination is resulting set from two condition hence 3 SELECT statements
- o column wise combination
- o Row wise combination
- o data type of two tables can be different datatype of corresponding columns from each table should be the same
- o can generate distinct or duplicate rows
- o the number of columns selected by the number of columns selected must be the same from each table
- o combines results horizontally
- o combines results vertically

customer

order-details



o UNION

- combines two or more SELECT statements

SELECT * FROM table1

UNION

- number of column, order of column must be same for table1 and table2

b INTERSECT

returns common values of the tables

Emulated

- * **SELECT DISTINCT column-list FROM table1 INNER JOIN table2 USING (join-column);**

c MINUS

- this operation returns the distinct row from the first table that does not occur in the second table

Emulated

* **SELECT column-list FROM table1 LEFT JOIN table2 ON condition WHERE table2.column-name IS NULL;**

d LEFT JOIN

e.g - **SELECT id FROM table1 LEFT JOIN table2 USING(id) WHERE table2.id IS NULL;**

full JOIN
SELECT c.*, o.* FROM customer AS c LEFT JOIN order-details AS o ON c.cust_id = o.order_id;

left JOIN
SELECT c.* FROM customer AS c LEFT JOIN order-details AS o ON c.cust_id = o.order_id;

right JOIN
SELECT o.* FROM customer AS c RIGHT JOIN order-details AS o ON c.cust_id = o.order_id;

cross JOIN
SELECT customer_id, order_details_id FROM customer CROSS JOIN

cross JOIN
SELECT customer_id, order_details_id FROM customer CROSS JOIN order_details;

Q. can we JOIN without JOIN keyword...?
→ Yes

SELECT * FROM customer, order-details WHERE customer.id = order-details.cust_id;

Advantages of JOIN
- better readability

$$A = \begin{pmatrix} 1 & 4 \\ 1 & 1 \end{pmatrix}, B = \begin{pmatrix} 1 & 0 \\ 1 & 2 \\ 1 & 3 \end{pmatrix}$$

$$A \cup B = \underline{\underline{1}}$$

UNION Select * From T₁ UNION T₂ where some columns are needed in T₁ & T₂

UNION

Select * From T₁

INTERSECTION SELECT DISTINCT id From T₁ INNER JOIN T₂ we use INNER JOIN for INTERSECTION

MINUS SELECT id From T₁ LEFT JOIN T₂ MINUS(id) WHERE T₂.id IS NULL;

DEPT 1 DEPT 2

empid	name	role	empid	name	role
1	A	engineer	3	C	manager
2	B	salesman	6	E	marketing
3	C	manager	7	G	salesman
4	D	salesman	8	H	engineer
5	E	engineer			

□ SUB QUERIES

- SELECT * FROM customer WHERE id IN (SELECT id FROM customer WHERE c_name = 'TAHAT');

- SELECT * FROM employee WHERE age IN (SELECT age FROM employee WHERE age >= 30);

Q Find employee details working in more than 1 project

- SELECT * FROM employee WHERE id IN (SELECT empID FROM project group by empID having count(empID) > 1);

Q select employee details having age > avg(age)

- SELECT * FROM employee WHERE age > (SELECT avg(age) FROM employee);

Q select max age person whose first name contains 'al'

- SELECT max(age) FROM (SELECT * FROM employee WHERE fname like 'al%')

As temp alias

new table name

- Q. list out all the employee in all department who work as salesman
- SELECT * FROM dept1 WHERE role = 'SALESMAN'

UNION

SELECT * FROM dept2 WHERE role = 'SALESMAN';

↳ union all the results of both the queries now & then

* find 3rd oldest employee

32
44
22
31

SELECT * FROM employee WHERE 3 = 6

SELECT COUNT(ee.age)

FROM employee ee
WHERE ee.age >= (SELECT ee.age
FROM employee ee
WHERE ee.age >= ee.age))

→ 2

ee.age = 44

→ 1

3) ee.age = 32
→ 4

ee.age = 32

→ 3

ee.age = 31
→ 3

ee.age = 31

→ 2

VIEW

CREATE VIEW customer_view AS SELECT fname, age FROM employee;

showed - SELECT * FROM customer_view;

ALTER VIEW customer_view AS SELECT fname FROM employee;

DROP VIEW IF EXISTS customer_view;

- Sub Queries**
- outer query depends upon inner query alternatives to John's
 - nested queries
 - SELECT column-list FROM table-name WHERE column-name OPERATOR (SELECT column-list FROM table-name [WHERE])
- eg - SELECT * FROM table1 WHERE col1 IN (SELECT col1 FROM table1);
- o sub queries exist mainly in 3 clause
 - 1. inside a WHERE clause
 - 2. inside a FROM clause
 - 3. inside a SELECT clause
- ↳ Subquery using FROM clause
- SELECT MAX(salary) FROM (SELECT * FROM movie WHERE country='India') as temp
- ↳ Subquery using SELECT
- SELECT (SELECT column-list FROM tablename WHERE condition), column-list FROM tablename WHERE condition.
 - Derived subquery
 - SELECT column-list FROM (SELECT column-list FROM tablename WHERE [condition]) as viewable_name;
- ↳ Co-related Sub-query
- with a normal nested subquery, the inner select query runs first and executes once, returning values to be used by the main query
 - correlated subquery, however, executes once for each candidate row considered by the outer query, in other words, the inner query is driven by the outer query.

JOINS

Faster

- Joins maximise calculation burden on DBMS

- complex, difficult to understand and implement
- choosing optimal join for optimal - easy use case is difficult.

MySQL VIEWS

- a view is database object that has no values. Its contents are based on the base table it contains row and columns similar to the real table
- in MySQL, the view is a virtual table created by a query by joining one or more tables. It is operated similarly to the base table but does not contain any data of its own.
- the view and table have one main difference that the view are definitions built on top of other tables (or views). If any changes occur in the underlying table, the same changes reflected in the view also.
- - CREATE VIEW view-name AS SELECT columns FROM table [WHERE conditions];
 - ALTER VIEW view-name AS SELECT columns FROM table WHERE conditions;
 - DROP VIEW IF EXISTS view-name;
 - CREATE VIEW trainer AS SELECT course_name, strahler, t_email FROM courses t, contact t WHERE t.id = t_id; Views using join clauses
- note : We can also import/export table schema from files (.csv or .json)

SUBQUERIES

Slower

- keeps responsibility of calculation on user

- comparatively easy to understand and implement

- choosing optimal join for optimal - easy use case is difficult.

MySQL VIEWS

- a view is database object that has no values. Its contents are based on the base table it contains row and columns similar to the real table
- in MySQL, the view is a virtual table created by a query by joining one or more tables. It is operated similarly to the base table but does not contain any data of its own.
- the view and table have one main difference that the view are definitions built on top of other tables (or views). If any changes occur in the underlying table, the same changes reflected in the view also.
- - CREATE VIEW view-name AS SELECT columns FROM table [WHERE conditions];
 - ALTER VIEW view-name AS SELECT columns FROM table WHERE conditions;
 - DROP VIEW IF EXISTS view-name;
 - CREATE VIEW trainer AS SELECT course_name, strahler, t_email FROM courses t, contact t WHERE t.id = t_id; Views using join clauses
- note : We can also import/export table schema from files (.csv or .json)

QUESTION AND ANSWERS

Slower

- keeps responsibility of calculation on user

- comparatively easy to understand and implement

- choosing optimal join for optimal - easy use case is difficult.

MySQL VIEWS

- a view is database object that has no values. Its contents are based on the base table it contains row and columns similar to the real table
- in MySQL, the view is a virtual table created by a query by joining one or more tables. It is operated similarly to the base table but does not contain any data of its own.
- the view and table have one main difference that the view are definitions built on top of other tables (or views). If any changes occur in the underlying table, the same changes reflected in the view also.
- - CREATE VIEW view-name AS SELECT columns FROM table [WHERE conditions];
 - ALTER VIEW view-name AS SELECT columns FROM table WHERE conditions;
 - DROP VIEW IF EXISTS view-name;
 - CREATE VIEW trainer AS SELECT course_name, strahler, t_email FROM courses t, contact t WHERE t.id = t_id; Views using join clauses
- note : We can also import/export table schema from files (.csv or .json)

Lecture 10 IMP 50 questions

Q. Worker

WORKER-ID	FIRST-NAME	LAST-NAME	SALARY	HIRING-DATE	DEPARTMENT
1	Manika	Aroa	100000	2014-02-20 09.00.00	HR
2	Nikharika	Verma	80000	2014-06-11 09.00.00	Admin
3	Vishal	Singhal	300000	2014-02-20 09.00.00	HR
4	Amilabh	Singh	500000	2014-02-20 09.00.00	Admin
5	Vineet	Bhati	500000	2014-06-11 09.00.00	Admin
6	Vipul	Divan	20000	2014-06-11 09.00.00	Account
7	Satish	Kumar	75000	2014-01-20 09.00.00	Account
8	Gretika	Chauhan	90000	2014-04-11 09.00.00	Admin

Q. Bonus

WORKER-REF-ID	BONUS-AMOUNT	BONUS-DATE
1	5000	2016-02-20 00.00.00
2	3000	2016-06-11 00.00.00
3	4000	2016-02-20 00.00.00
1	4500	2016-02-20 00.00.00
2	3500	2016-06-11 00.00.00

Q. 5 write query to find the position of the alphabet ('b') in the first name column
 → SELECT instr(FIRST-NAME, 'b') FROM worker WHERE FIRST-NAME = 'Amilabh';

Q. 6 write query to print the FIRST-NAME from worker table after removing white spaces from the right side
 → SELECT trim(FIRST-NAME) FROM worker;

Q. 7 write query to print the DEPARTMENT from worker table after removing white spaces from the left side
 → SELECT ltrim(FIRST-NAME) FROM worker;

Q. 8 write query that fetches the unique values of DEPARTMENT from worker table and prints its length.
 → SELECT DISTINCT DEPARTMENT, length(DEPARTMENT) FROM worker;

Q. 9 write query to print the FIRST-NAME from worker after replacing 'a' with 'A'
 → SELECT replace(FIRST-NAME, 'a', 'A') FROM worker;

Q. 10 write an SQL query to fetch "FIRST-NAME" from worker table using the alias name as <WORKER_NAME>.

→ SELECT FIRST-NAME AS WORKER_NAME FROM worker;

Q. 2 write query to fetch "FIRST-NAME" from worker table in upper case
 → SELECT uppercase(FIRST-NAME) FROM worker;

Q. 3 write query to fetch unique values of DEPARTMENT from worker table
 → SELECT DISTINCT DEPARTMENT FROM worker;

Q. 4 write query to print first three characters of FIRST-NAME from worker table
 → SELECT substr(FIRST-NAME, 1, 3) FROM worker;

- Q. 10 write query to print the FIRST-NAME and LAST-NAME from worker table into a single column CONCATE-NAME AS CONCATE-NAME FROM worker;
- SELECT * FROM worker ORDER BY FIRST-NAME;
- Q. 11 write query to print all worker details from the worker table order by FIRST-NAME Ascending.
- SELECT * FROM worker ORDER BY FIRST-NAME;
- Q. 12 write query to print all worker details from the worker table order by FIRST-NAME ascending and DEPARTMENT descending.
- SELECT * FROM worker ORDER BY FIRST-NAME, DEPARTMENT DESC;
- Q. 13 write query to print details of workers with the FIRST-NAME as "Vipul" and "Sathish" from worker table.
- SELECT * FROM worker WHERE FIRST-NAME IN ('Vipul', 'Sathish');
- Q. 14 write query to print details of workers excluding FIRST-NAME "Vipul" and "Sathish" from worker table
- SELECT * FROM worker WHERE FIRST-NAME NOT IN ('Vipul', 'Sathish');
- SELECT * FROM worker WHERE FIRST-NAME <> 'Vipul' AND FIRST-NAME <> 'Sathish';
- Q. 15 write query to print details of workers with department name as "Admin"
- SELECT * FROM worker WHERE DEPARTMENT LIKE 'Admin%';
- Q. 16 write query to print details of the worker whose FIRST-NAME contains 'a'
- SELECT * FROM worker WHERE FIRST-NAME LIKE '%a%';
- Q. 17 write query to print details of the workers whose FIRST-NAME ends with 'a'.
- SELECT * FROM worker WHERE FIRST-NAME LIKE '%a';
- Q. 18 write query to print details of the workers whose first name ends with 'h' and contains 6 alphabets
- SELECT * FROM worker WHERE FIRST-NAME LIKE '----h';

- Q. 19 write query to print details of the workers whose salary lies between 1.00,000 AND 5,00,000.
- SELECT * FROM worker WHERE SALARY BETWEEN 100,000 AND 5,00,000;
- Q. 20 write query to print details of the workers who have joined in Feb '2014.
- SELECT COUNT(DISTINCT FIRST-NAME, LAST-NAME) FROM worker WHERE YEAR(joining-date) = 2014 AND MONTH(joining-date) = 02;
- Q. 21 write query to fetch the count of employee working in the department 'Admin'
- SELECT COUNT(DISTINCT DEPARTMENT) FROM worker WHERE DEPARTMENT = 'Admin';
- Q. 22 write query to fetch worker full names with salaries \geq 50,000 and \leq 1,00,000;
- SELECT COUNT(FIRST-NAME, LAST-NAME) FROM worker WHERE SALARY BETWEEN 50,000 AND 1,00,000;
- Q. 23 write query to fetch the no. of workers for each department in the descending order.
- SELECT DEPARTMENT, COUNT(WORKER-ID) FROM worker GROUP BY DEPARTMENT ORDER BY COUNT(WORKER-ID) DESC;
- Q. 24 write query to print details of the worker who are also manager.
- SELECT w.* FROM worker AS w INNER JOIN title AS t ON worker-id = worker-ref-id WHERE worker-title = 'Manager';
- Q. 25 write query to fetch number (more than 1) of same titles in the imp of different types.
- SELECT WORKER-TITLE, COUNT(WORKER-TITLE) FROM title GROUP BY WORKER-TITLE HAVING COUNT(WORKER-TITLE) > 1;

Q. 26 write query to show only odd rows from a table

→ SELECT * FROM worker WHERE MOD(WORKER-ID, 2) != 0;

Q. 27 write query to show only even rows from a table

→ SELECT * FROM worker WHERE MOD(WORKER-ID, 2) = 0;

Q. 28 write query to clone a new table from another table

→ CREATE TABLE worker-done LIKE worker;

INSERT INTO worker-done SELECT * FROM worker;

SELECT * FROM worker-done;

Q. 29 write query to fetch intersecting records of two tables.

→ SELECT worker.* FROM worker INNER JOIN worker-done USING(WORKER-ID);

Q. 30 write query to show records from one table that another table does not have

→ SELECT worker.* FROM worker LEFT JOIN worker-done using (WORKER-ID)
WHERE worker-done.WORKER-ID IS NULL;

minus

Q. 31 write query to show the current date and time.

→ SELECT curdate();

SELECT now();

Q. 32 Write query to show the top n (say 5) records of a table ordered by decreasing salary

→ SELECT * FROM worker ORDER BY SALARY DESC LIMIT 5;

Q. 33. write query to determine the nth (say n=5) highest salary from a table

→ SELECT * FROM worker ORDER BY SALARY DESC LIMIT 4,1;

4
5
3
6
salary
sooro
sooco
vick
300000
200000
Vimal
means first 4 not take from 4 to further 2 take

means 1 take

Q. 34 write query to determine the 5th highest salary without using LIMIT keyword

→ SELECT WORKER-ID, SALARY FROM worker AS w1 WHERE n <= (

SELECT count(DISTINCT(w2.salary)) FROM worker AS w2

WHERE w2.salary >= w1.salary);

Q. 35 write query to fetch the list of employee with the same salary

→ SELECT w1.* FROM worker AS w1, worker AS w2 USING(Salary)

WHERE w1.WORKER-ID <> w2.WORKER-ID;

→ SELECT w1.* FROM worker AS w1, worker AS w2 WHERE

w1.salary = w2.salary AND w1.WORKER-ID != w2.WORKER-ID;

Q. 36 write query to show second highest salary from table using sub-query

→ SELECT max(salary) FROM worker
WHERE salary NOT IN (SELECT max(salary) FROM worker);

Q. 37 write query to show one row twice in result from a table

→ SELECT * FROM worker UNION ALL SELECT * FROM worker ORDER BY worker-ID;

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

union ALL → common how many occurred in both

Q. 38 write query to list worker-ID who does not get bonus.

→ SELECT WORKER-ID FROM worker WHERE WORKER-ID NOT IN (

SELECT WORKER-ID FROM bonus);

Q. 39 write query to fetch the first 50% records from a table.

→ SELECT * FROM worker WHERE WORKER-ID <= (SELECT count(WORKER-ID)/2 FROM

worker);

- Q. 40 write query to fetch the departments that have less than 4 people in it.
- `SELECT department, count(department) FROM worker GROUP BY department HAVING COUNT(department) < 4;`
- Q. 41 write query to show all departments along with the number of people there.
- `SELECT department, count(department) FROM worker GROUP BY department;`
- Q. 42 write query to show the last record from a table.
- `SELECT * FROM worker WHERE worker_id = (SELECT max(worker_id) FROM worker);`
- Q. 43 write query to fetch the first row of a table
- `SELECT * FROM worker WHERE worker_id = (SELECT min(worker_id) FROM worker);`
- Q. 44 write query to fetch the last five records from a table
- `(SELECT * FROM worker ORDER BY worker_id DESC LIMIT 5) ORDER BY worker_id;`
- Q. 45 write query to print the name of employee having the highest salary in each department.
- `EFFECT w.department, w.first_name, w.salary FROM (SELECT max(salary) AS maxsal, department FROM worker GROUP BY department) temp INNER JOIN worker w ON temp.department = w.department AND temp.maxsal = w.salary;`
- * Q. 46 write query to fetch three max salaries from a table using correlated subquery
- `SELECT distinct salary FROM worker AS w1 WHERE 3 >= (SELECT count(distinct salary) FROM worker AS w2 WHERE w1.salary <= w2.salary ORDER BY w2.salary DESC);`
- `SELECT DISTINCT salary FROM worker ORDER BY salary DESC LIMIT 3;`

- Q. 47 write query to fetch three min salaries from a table using correlated subquery
- `SELECT distinct salary FROM worker w1 WHERE 3 >= (SELECT count(distinct salary) FROM worker w2 WHERE w1.salary >= w2.salary ORDER BY w2.salary DESC);`
- Q. 48 write query to fetch the n^{th} max salaries from a table
- `SELECT distinct salary FROM worker w1 WHERE n >= (SELECT count(distinct salary) FROM worker w2 WHERE w1.salary <= w2.salary) ORDER BY w2.salary DESC;`
- Q. 49 write query to fetch departments along with the total salaries paid for each of them
- `SELECT department, sum(salary) FROM worker GROUP BY department ORDER BY sum(salary);`
- Q. 50 write query to fetch the name of workers who earn the highest salary in each department.
- `SELECT first_name, salary FROM worker WHERE salary = (SELECT max(salary) FROM worker);`

last question

A Pairs

A	B	A	B	A	B
1	2	1	2	2	1
2	4	2	4	2	1
1	5	6	5	6	5
3	2	3	2	3	2
4	2	4	2	4	2
5	6	5	6	1	2
6	5	2	4	3	2
7	8	7	8	5	6
8	7	6	5	7	8

Q.1

remove reversed pairs JOINs

→ SELECT t1.* FROM pairs AS t1 LEFT JOIN pairs AS t2 ON t1.A = t2.B AND t1.B = t2.A WHERE t1.A IS NULL OR t1.A < t2.A;

Q.2 print only that pairs which are not repeated

→ SELECT t1.* FROM pairs AS t1 LEFT JOIN pairs AS t2 ON t1.A = t2.B AND t1.B = t2.A WHERE t1.A IS NULL AND t1.B IS NULL;

Q.2 remove reversed pairs (correlated subquery)

→ SELECT * FROM pairs AS p1 WHERE NOT EXISTS (SELECT * FROM pairs AS p2 WHERE p1.B = p2.A AND p1.A > p2.B AND p1.A > p2.B);

- normalisation is a step towards DB optimisation

a) Functional Dependency FD

- it's relationship between the primary key attribute (usually) of the relation to that of the other attribute of the relation

$X \rightarrow Y$ left side $X =$ Determinant, right side $Y =$ Dependent

a) Types of FD

① Trivial FD : $A \rightarrow B$ has trivial functional dependency if B is a subset of A

$A \rightarrow A, B \rightarrow B$ are also trivial FD

② Non-trivial FD : $A \rightarrow B$ has a non-trivial functional dependency if B is not a subset of A . (A intersection B is null)

a) Rules of FD Armstrong's axioms

- Reflexive
- if ' A' is a set of attributes and ' B' ' is subset of ' A ', then $A \rightarrow B$ holds
- if $A \supseteq B$ then $A \rightarrow B$

b) Augmentation

- if B can be determined by A , then adding an attribute to this functional dependency won't change anything

- if $A \rightarrow B$ holds, then $A \cup X \rightarrow B \cup X$ holds too. ' X ' being a set of attributes

c) Transitivity

- if A determines B and B determines C , we can say that A determines C .

- if $A \rightarrow B$ and $B \rightarrow C$ then $A \rightarrow C$

- o Why normalisation
 - to avoid redundancy in DB, not to store redundant data.

o what happens if we have redundant data

- insertion, deletion, update anomalies arises

■ Anomalies

- ⇒ Insertion anomaly
when certain data (attribute) can not be inserted into the DB without the presence of other data

2) Deletion anomaly

- The deletion anomaly refers to that situation where the deletion of data results in the unintended loss of some other important data.

3) Update anomaly or modification anomaly

- The update anomaly is when an update of a single data value requires multiple rows of data to be updated.
- due to update to many places, may be data inconsistency arises, if one forgets to update the data at all the intended places.

a) 3NF

- relation must be in 2NF.
- no transitivity dependency exists

b) BCNF boyce - codd normal form

- relation must be in 3NF
- FD: A \rightarrow B, A must be a super key

- we must not derive prime attribute from any prime or non-prime attribute

Q What is normalisation

- normalisation is used to minimize the redundancy from a relation; it is also used to eliminate undesirable characteristics like insertion, update, deletion anomalies
- normalisation divides the composite attributes into individual attributes, or larger table into smaller and links them using relationship
- the normal form is used to reduce redundancy from the database table.

a) Advantages of normalisation

- normalisation helps to minimize data redundancy
- greater overall database organization
- data consistency is maintained in DB

■ Types of normal forms

INF, 2NF, 3NF, BCNF

a) INF

- every relation cell must have atomic value
- relation must not have multi-valued attributes

b) 2NF

- relation must be in INF
- there should not be any partial dependency
- all non-prime attributes must be fully dependent on PK
- non-prime attributes cannot depend on the part of the PK

c) 3NF

- relation must be in 2NF
- no transitivity dependency exists
- non-prime attributes should not find a non-prime attributes

d) BCNF

- relation must be in 3NF
- FD: A \rightarrow B, A must be a super key
- we must not derive prime attribute from any prime or non-prime attribute

Transaction

comes forward to DB

1) Transaction

- A unit of work done against the DB in a logical sequence.
- sequence is very important in transaction.
- it is a logical unit of work that contains one or more SQL statements.
- the result in all this statements in a transaction either gets completed successfully (all the changes made to the DB are permanent) or if at any point any failure happens, it gets rollbacked (all the changes being done are undone).

read(A)

A = A - 50

write(A)

B = B + 50

read(B)

write(B)

while(B)

if B > 0

 B = B - 1

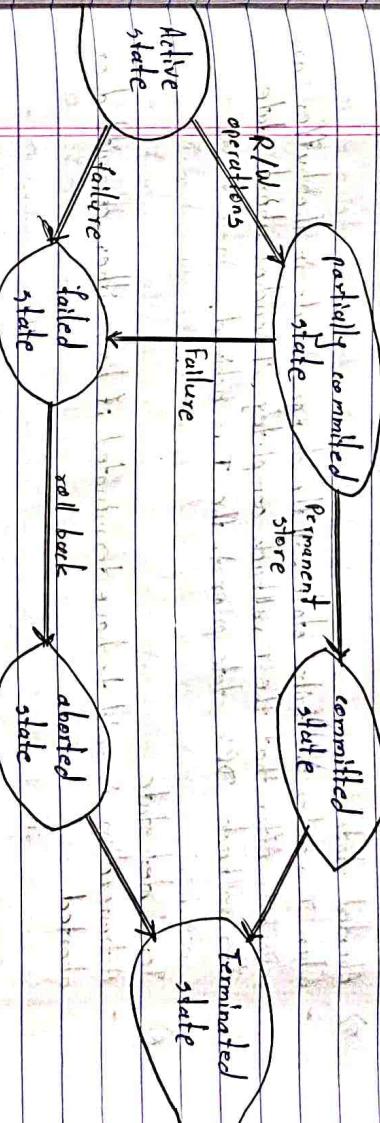
end if

end while

end transaction

2) Durability

after transaction completes successfully, the changes it had made to the database persist, even if there are system failure.



3) Atomicity

- either all operations of transaction are reflected properly in the DB, or none.
- to ensure consistency of the data, we require that the DB system maintains the following properties of the transaction.

4) Atomicity

- either all operations of transaction are reflected properly in the DB, or none.

ACID properties

- to ensure consistency of the data, we require that the DB system maintains the following properties of the transaction.

1) Atomicity

- either all operations of transaction are reflected properly in the DB, or none.

5) Isolation

- even though multiple transaction may execute correctly, the system guarantees that, for every pair of transaction T_i and T_j , if appears to T_i that either T_j finished execution before T_i started, or T_i started execution after T_j finished. Thus, each transaction is unaware of other transaction executing concurrently in the system.

- multiple transaction can happen in the system in isolation, without interfering each other.

6) Consistency

- when updates are made permanent on the DB, the the T is said to be in committed state. rollback can't be done from the commit state, now consistent state is achieved at this stage.

- even though multiple transaction may execute correctly, the system

- guarantees that, for every pair of transaction T_i and T_j , if appears to T_i that either T_j finished execution before T_i started, or T_i started execution after T_j finished. Thus, each transaction is unaware of other transaction executing concurrently in the system.

- multiple transaction can happen in the system in isolation, without interfering each other.

How to implement atomicity and durability in transactions

- 4) Failed state
- when T is being executed and some failure occurs, due to this, it is impossible to continue the execution of the T.

- 5) Aborted state
- when T reaches the failed state, all the changes made in the buffer are reversed. After that the T rollback completely. T reaches abort state after rollback. DB's state prior to the T is achieved.

6) Terminated state

- A transaction is said to have terminated if has either committed or aborted.

- o if T success, it is committed as,

- DB makes sure all the pages of the new copy of DB written on the disk
- DB system updates the db-pointer to point to the new copy of DB
- new copy is now the current copy of DB
- the old copy is deleted
- the T is said to have been COMMITTED at the point where the updated db-pointer is written to disk

7) Atomicity

- if T fails at any time before db-pointer is updated, the old content of DB are not affected

- T abort can be done by just deleting the new copy of DB
- Hence, either all updates are reflected or none

8) Durability

- suppose, system fails at any time before the updated db-pointer is written to disk
- when the system restarts, it will read db-pointer and will thus, see the original content of DB and none of the effects of T will be visible
- T is assumed to be successful only when db-pointer is updated

- Recovery mechanism component of DBMS supports atomicity and durability
- shadow-copy schema
- based on making copies of DB aka shadow copies
- assumption only one transaction (T) is active at a time
- a pointer called DB pointer is maintained on the disk, which at any instance points to current copy of DB
- T that wants to update DB first create a complete copy of DB.
- all further updates are done on new DB copy leaving the original copy (shadow copy) untouched.
- if at any point the T has to be aborted the system deletes the new copy and the old copy is not affected

- if system fails after db-pointer has been updated before that all the pages of the new copy were written to disk. hence, when system restarts, it will read new db copy

the implementation is dependent on write to the db-pointer being atomic.

Luckily, disk system provide atomic updates to entire block or at least a disk sector. So we make sure db painter lies entirely in a single sector.

By sorting db-pointer at the beginning of a block, entire DB is copied for every transaction.

卷之三

10/01/11 00-2011-11

old JB copy

new copy (RAM) - control main memory

new copy on **DISK**

all operations on
new copy

卷之三

log-based recovery methods

The log is a sequence of records. Log of each transaction is maintained in stable storage so that any failure occurs then it can be recovered.

If any operation is performed on the database, then it will be

the process of sorting the logs should be done before they are applied in the database.

stable storage is a classification of computer data storage.

be written that is robust against some hardware anomalies occurring for any given write operation and also

LAW OF THE STATE OF CALIFORNIA

 b $\hat{=}$ 2000	logs (stable storage)	Immediate DB modifications.
 A $\hat{=}$ 1000	logs (stable storage)	Immediate DB modifications.

a) Deferred DB modifications

- ensuring atomicity by recording all the DB modification in the log but determining the execution of all the write operation until the final action of the T has been executed
 - log information is used to executing deferred writes when T is completed
 - if system crashed before the T completes, or if T aborted, the information in the logs are ignored
 - if T completes, the records associated to it in the log file are used in executing the ~~execution of all the write operations~~, deferred writes
 - if failure occurs while this updating is taking place, we perform redo

■ Immediate DB modifications

- DB modifications to be output to the DB while the T is still in active state
 - DB modifications written by active T are called uncommitted modifications
 - in the event of crash or T failure, system uses old value field of the log records to restore modified values.
 - update takes place only after log records in a stable storage

Failure handling

- system failure before T completes, or if T aborted, then old value is used to undo the T.
 - if T completes and system crashes, then new value field is used to redo T having commit logs on the logs

- Indexing is used to optimise the performance of database by minimising the number of disk access required when a query is processed.
 - the index is a type of data structure. it is used to locate and access the data in database table quickly.
 - speeds up operation with read operations like SELECT query , WHERE clause etc.
 - Search key : contains copy of primary key or candidate key of the table or something else.
 - Data Reference : pointer holding the address of disk block where the value of the corresponding key is stored.
 - indexing is optional, but increases access speed. it is not the primary mean to access the tuple. it is the secondary mean.
- o Index file is always sorted.**
- | search key | data reference | value |
|------------|----------------|-------|
| key | key | key |
- ④ Indexing methods.**
- ① Primary Index - clustering index**
- A file may have several indices. On different search keys. if the data file containing the records is sequentially ordered, a primary index is an index whose search key also defines the sequential order of the file.
 - Note** - the term primary index is sometimes used to mean an index on primary key. However such usage is non-standard and should be avoided.
 - all files are ordered sequentially on some search key. It could be primary key or non-primary key.
- o Dense Index**
- the dense index contains an index record for every search key value in the file.
 - the index record contains the search key value and a pointer to the first data record with that search key value. the rest of the records with the same search key value would be stored sequentially after the first record.
 - it needs more space to store index record itself. the index records have the search key and a pointer to the actual record on the disk.

pr key

address

1	0x001	1
2	0x002	2
3	0x003	3
4	0x004	4
5	0x005	5

data sheet index

data base

Sparse Index

- an index record appears for only some of the search-key values.
- sparse index helps you to resolve the issues of dense indexing in DBMS. In this method of indexing technique, a range of index columns stores the same data block address, and when data needs to be retrieved, the block address will be fetched.
- primary indexing can be based on data file is sorted w.r.t primary key attribute or non-key attributes.

- o Based on key attribute
 - data file is sorted wrt primary key attribute.
 - data file is sorted wrt primary key attribute value in the data file.
 - this is denoted as, all the unique values have an entry in the index file.

eg - let's assume that a company recruited many employees in various departments. In this case, clustering indexing in DBMS should be created for all employees who belong to the same dept.

index stores first occurrence in DB

pr key

Block

1	B1	1
11	B2	2
21	B3	3
31	B4	4
41		5
51		6
61		7
71		8
81		9
91		10

Index

1	B1	1
2	B1	2
3	B1	3
4	B2	3
5	B3	3
6	B3	4
7	B3	4
8	B3	4
9	B3	5
10	B3	5
11	B3	6
12	B3	6
13	B3	7
14	B3	7
15	B3	8
16	B3	8
17	B3	9
18	B3	9
19	B3	10
20	B3	10
21	B3	10
22	B3	10
23	B3	10
24	B3	10
25	B3	10
26	B3	10
27	B3	10
28	B3	10
29	B3	10
30	B3	10
31	B3	10

N

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

19

20

21

22

23

24

25

26

27

28

29

30

31

32

33

34

35

36

37

38

39

40

B4

No SQL

- No SQL database (aka "not only SQL") are not tabular database and store data differently than relational table. NoSQL databases come in a variety of types based on their data model. The main types are document, key-value, wide-column and graph. They provide flexible schema and scale easily with large amounts of data and high user loads.
- they are schema free
- data structures used are not scalar. They are more flexible, has the ability to adjust dynamically
- can handle huge amount of data (big data)
- most of the NoSQL are open sources and has the capability of horizontal scaling if it just stores data in some format other than relational.

History behind No SQL

- NoSQL databases emerged in the last 2000s as the cost of storage dramatically decreased, gone were the days of needing to create a complex, difficult-to-manage data model in order to avoid data duplication. Developers (rather than storage) were becoming the primary cost of software development, so NoSQL databases optimized for developer productivity.
- data becoming unstructured more, hence structuring (defining schema in advance) them had becoming costly
- NoSQL allows developer to store huge amount of unstructured data, giving them a lot of flexibility
- recognizing the need to rapidly adapt to changing requirements in a software system, developers needed the ability to iterate quickly and make changes throughout their software stack - all the way down to the database. NoSQL databases gave them their flexibility.
- Cloud computing also rose in popularity, and developers began using public clouds to host their applications and data. They wanted the ability to distribute data across multiple servers and regions to make their applications resilient, to scale out instead of scale up, and to intelligently geo-place their data. Some NoSQL databases like mongoDB provide these capabilities.

No-SQL databases Advantages.

- ① Flexible Schema**
 - RDBMS has pre-defined schema, which become an issue when we do not have all the data with us or we need to change the schema. It's a huge task to change schema on the go.

② Horizontal Scaling

- horizontal scaling also known as scale-out refers to bringing on additional nodes to share the load. This is difficult with relational databases due to the difficulty in spreading out related data across nodes, with non-related DB's, this is made simpler since collections are self-contained and not coupled relationally. This allows them to be distributed across nodes more simply, as queries do not have to "join" them together across nodes.
- scaling horizontally is achieved through sharding OR replica sets

③ High Availability

- NoSQL DB's are highly available due to its auto replication feature, i.e. whenever any kind of failure happens, data replicates itself to the preceding consistent state.
- If a server fails, we can access that data from another server as well, as in NoSQL database data is stored at multiple servers.

④ Easy Insert & Read operations

- queries in NoSQL databases can be faster than SQL database, why? data in SQL databases is typically normalized, so queries for a single object or entity require you to join data from multiple tables. As your table grows in size, the joins can become expensive. However, data in NoSQL database is typically stored in a way that is optimized for queries, the rule of thumb when you use MongoDB is data that is accessed together should be stored together, queries typically do not have or require joins, so the queries are very fast.

- but difficult to delete or update operations.
- Caching mechanism
- No-SQL use case is more for cloud applications

When to use No-SQL?

- fast-paced agile development
- storage of structured and semi-structured data
- huge volume of data
- requirements for scale-out architecture
- modern application paradigms like micro-services & real-time streaming

No-SQL DB Misconceptions

- relationship data is best suited for relational databases
- A common misconception is that No-SQL databases or non-relational databases don't store relationship data well. No-SQL databases can store relationship data. - they just store it differently than relational databases do. in fact, when compared with relational databases, because related data doesn't have to split between tables. No-SQL data models allow nested within a single data structure.
- No-SQL databases don't support ACID transactions.
- another common misconception is that No-SQL databases don't support ACID transactions. some No-SQL databases like mongoDB do, in fact, support ACID transactions.

No-SQL DB Models

- ① Key-value stores
 - the simplest type of No-SQL database is a key-value store. every data element in the database is stored as a key-value pair consisting of an attribute name (or "key") and a value. in a sense, a key-value store is like a relational database with only two columns: the key or attribute name

(such as "state") and the value (such as "Alaska")

use cases include shopping carts, user preferences, and user profiles.

Eg. Oracle No-SQL, Amazon DynamoDB, Mongo-DB also supports key-value store.redis.

a key-value associates a value (which can be anything from a number or simple string to a complex object) with a key, which is used to keep track of the object. in its simplest form, a key-value store is like a dictionary/array/map object as it exists in most programming paradigms, but which is stored in a persistent way and managed by a database management system. key-value databases use compact, efficient index structures to be able to quickly and reliably locate a value by its key, making them ideal for systems that need to be able to send an receive data in consistent time.

there are several usecases where choosing a key-value store approach is an optimal solution.

- ② real-time random data access, e.g. - user session attributes in an online application such as gambling or finance
- ③ caching mechanism from frequently accessed data or configuration based on key
- ④ application is designed on simple key-based queries

② Column-Oriented / Columnar / C-store / Wide-column:

the data is stored such that each row of the column will be next to other rows from that same column.

while a relational database stores data in rows and reads data row by row, a column store is organised as a set of columns. this means, that when you want to run analytics on a small number of columns, you can read those columns directly without consuming memory with the unwanted data. columns are often of the same type and benefit from more efficient compression, making reads even faster. columnar databases can quickly aggregate the value of a given column (adding up the total sale for the year, for example).

use cases include analytics

- cassandra, redshift, snowflake.

③ Document-based storage

- This DB stores data in documents similar to JSON (JavaScript object notation) objects. each document contains pairs of fields and values. The values can typically be a variety of types, including strings, numbers, boolean, arrays, or objects.
- use cases include e-commerce platforms, trading platform, and mobile app development across industries.
- supports ACID properties hence, suitable for transactions
- Eg - MongoDB, Couch DB.

④ Graph-based storage

- a graph database focuses on the relationship between data elements. each element is stored as a node (such as a person in a social media graph). The connections between elements are called links or relationships. In a graph database, connections are first-class elements of the database, stored directly. in relational databases, links are implied, using data to express the relationships.
- a graph database is optimized to capture and search the connections between data elements, overcoming the overhead associated with JOINing multiple tables in SQL
- very few real-world business systems can survive solely on graph queries. As a result, graph databases are usually run alongside other more traditional databases.
- use cases include fraud detection, social networks, and knowledge graphs.

	SQL DB	No-SQL DB
Data storage model	Tables with fixed rows and columns,	Document : JSON documents, Key-value : key-value pairs, Wide-column : tables with rows & dynamic columns Graph : nodes & edges
Difference Between SQL & No-SQL	Development history : development in the 1920's with a focus on reducing data duplication. A focus on reducing data scaling and allowing for rapid application change driven by agile & DevOps practices.	

Examples	Oracle, MySQL, Microsoft SQL Server, PostgreSQL	Document : MongoDB & couchDB key-value : Redis & DynamoDB Wide-column : Cassandra & HBase Graph : Neo4j & Amazon Neptune
Scaling	Vertical (scale-up)	Horizontal (scale-out across commodity servers)

Primary purpose	General purpose
Document	general purpose
Key-value	large amounts of data with simple look up queries
Wide column	large amounts of data with predictable query patterns.
Graph	analyzing & traversing relationships between connected data.
ACID properties	Supported
JOINS	Typically required
Data to object mapping	Typically not required
MongoDB documents map directly to data structures in most popular programming languages	many do not require ORMs MongoDB documents map directly to data structures in most popular programming languages

- ① Relational databases
- based on relational model
 - relational databases are quite popular, even though it was a system designed in the 1970s also, also known as relational database management systems (RDBMS), relational databases commonly use structured query language (SQL) for operations such as creating, reading, updating, deleting data. relational databases store information in discrete tables, which can be JOINed together by fields known as foreign keys. for example, you might have a user table which contains information about all your users, and join it to a purchases table, which contains information about all the purchases they've made. MySQL, Microsoft SQL server, and Oracle are types of relational DB's that are ubiquitous, having acquired a steady user base since the 1970s.
 - they provide a stronger guarantee at data normalisation
 - they use a well-known querying language through SQL
 - scalability issues (horizontal scaling)
 - data become huge, system become more complex

②

- Object oriented databases
- the object-oriented data model is based on the object-oriented-programming paradigm, which is now in wide use. Inheritance, Object-identity and Encapsulation (information hiding), with methods to provide an interface to objects, are among the key concepts of object-oriented programming that have found applications in data modeling. the object-oriented data model also supports a rich type system, including structured and collection types, while inheritance and object-identity distinguish the object-oriented data model from the E-R model
 - sometimes the database can be very complex, having multiple relations, so maintaining a relationship between them can be tedious at times,

- In object-oriented DB's data is treated as an object
- all bits of information come in one instantly available object package instead of multiple tables

④ Advantages

- data storage and retrieval is easy and quick
- can handle complex data relations and more variety of data types than standard relational databases.
- relatively friendly to model the advance real world problems.
- works with functionality of SQLs and Object Oriented languages

④ Dis-advantages

- High complexity causes performance issues like read, write, update and delete operations are slowed down.
- Not much of a community support as isn't widely adopted as relational databases
- Does not support views like relational databases.

Eg - Object DB, GemStone etc.

⑤ NO-SQL databases

- NO-SQL databases are non-tabular database and store data differently than relational tables. NO-SQL databases come in variety of types, based on their data model. The main types are document, key-value, wide-column, and graph. They provide flexible schema and scale easily with large amounts of data and high user loads.
- they are schema free
- data structures used are not tabular, they are more flexible than the ability to adjust dynamically
- can handle huge amount of data. (big data)
- most of the NO-SQL are open-sources and has the capability of horizontal scaling

⑥ Hierarchical database

- as the name suggests, the hierarchical database is most appropriate for use cases in which the main focus of information gathering is based on a concentric hierarchy, such as several individual employees reporting to a single department at a company
- the schema for hierarchical database is defined by its tree-like organisation, in which there is typically a root "parent" directory of data stored as records that links to various other subdirectory branches, and each subdirectory branch or child record, may link to various other subdirectory branches
- the hierarchical database structure dictates that, while a parent record can have several child records, each child record can only have one parent record. data within records is stored in the form of fields, and each field can only contain one value. retrieving hierarchical data from a hierarchical database architecture requires traversing the entire tree, starting at the root node.
- since the disk storage system is also inherently a hierarchical structure, these models can also be used as physical models.
- the key advantage of a hierarchical database is its ease of use. the one-to-many organisation of data makes traversing the database simple and fast, which is ideal for use cases such as website drop-down menus or computer folders in systems like Microsoft Windows OS. due to the separation of the tables from physical storage structures, information can easily be added or deleted without affecting the entirety of the database. and most major programming languages offer functionality for reading tree structure databases.
- the major disadvantage of hierarchical databases is their inflexible nature. the one-to-many structure is not ideal for complex structures as it cannot describe relationships in which each child node has multiple parents, nodes, also the tree-like organisation of data requires top-to-bottom sequential searching which is time consuming and require repetitive storage of data in multiple different entities, which can be redundant

Clustering in DBM

- database clustering (making replica sets) is the process of combining more than one server or instance connecting a single database. Some time one server may not be adequate to manage the amount of data or the number of requests, that is when a data cluster is needed. DB clustering, SQL server clustering, and SQL clustering are closely associated with SQL is the language used to manage the databases information.
 - replicate the same data set on different server

Clustering in SQL database clustering (making replica-sets) is the process of combining more than one server or instance connecting a single database. Some time one server may not be adequate to manage the amount of data or the number of requests, that is, when a data cluster is needed. DB clustering, SQL server clustering, and SQL clustering are closely associated with SQL is the language used to manage the databases information.

replicate the same data set on different server.

Advantages

Data Redundancy - clustering with databases helps with data redundancy, as we store the same data at multiple servers. don't confuse this data redundancy repetition of the same data that might lead to some anomalies. The redundancy that clustering offers is required and is quite certain due to the synchronisation, in case any of the servers had to face a failure due to any possible reason, the data is available at other servers for access.

2) Load balancing -

or scalability doesn't come by default with the database. it has to be brought by clustering regularly. it also depends on the setup. basically, what load balancing does is allocating the workload among the different servers that are part of cluster. this indicates that more users can be supported and if for some reasons if a huge spike in the traffic appears, there is a higher assurance that it will be able to support the new traffic. one machine is not going to all get all the hits. this can provide scaling seamlessly as required. this links directly to high availability. without load balancing, a particular machine could get overworked and traffic would slow down, leading to decrement of the traffic to zero.

Partitioning & Sharding in DBMS - DB optimisation

3) **High availability**
 when you can access a DB, it implies that it is available. High availability refers the amount of time a database is considered available. The amount of availability you need greatly depends on the number of transactions you are running on your database and how often you are running any kind of analytics on your data. With database clustering, we can reach extremely high levels of availability due to load balancing and have extra machines. In case a server goes down, the databases will, however, be available.

A big problem can solve easily when it is chopped into several smaller sub-problems, that is what the partitioning technique does. It divides a big database containing data, materials, matrix and indexes, into smaller and barely slices of data called partitions. The partitioned tables are directly used by SQL queries without any alteration, once the database partitioned. The data definition language easily work on the smaller partitioned slices instead of handling giant database altogether. This is how partitioning cuts down the problems in managing large DB problem table.

Partitioning

It is the technique used to divide stored database objects into separate servers due to this, there is an increase in performance, controllability of data. We can manage huge chunks of data optimally, when we horizontally scale our machines/servers. We know that it gives us a challenging time dealing with relational databases, as it's quite tough to maintain the relations, but if we apply partitioning to the database that is already scaled out. If we equipped with multiple servers, we can partition our database among those servers and handle the big data easily.

Vertical Partitioning

Slicing relation vertically / column-wise
 - need to access different servers to get complete tuples.

Horizontal Partitioning

- slicing relation horizontally / row-wise
- independent chunks of data tuples are stored in different servers.

When Partitioning is Applied

- data become much huge that managing and dealing with it became a tedious task
- the number of requests are enough larger than the single DB server access. It is taking huge time and hence the system's response time become high.

Database Scaling Patterns

- Advantages of partitioning
 - parallelism
 - Availability
 - Performance
 - Manageability
 - Reduce cost, as scaling-up or vertical scaling might be costly

- A case study - cab booking App
 - Tiny startup
 - ~10 customer onboard
 - A single small machine DB stores all customers, trips, locations, booking data, and customer trip history
 - ~1 trip booking in 5 mins

- Distributed database
 - A single logical database that is spread across multiple locations (Servers) and logically interconnected by network
 - This is the product of applying DB optimisation techniques like Clustering, Partitioning, Sharding

Sharding

- Technique to implement horizontal partitioning
- the fundamental idea of sharding is the idea that instead of having all the data sit on one DB instance, we split it up and introduce a routing layer so that we can forward the request to the right instances that actually contain the data.

1.4 Query Optimisation & Connection Pool Implementation

- Pros
 - Scalability
 - Availability
 - Cons
 - complexity, making partition mapping, routing layer to be implemented in the system, non-uniformity that creates the necessity of Re-sharding
 - Not well suited for analytical type of queries, as the data is spread across different DB instances (Scatter-Gather problem)
- Solutions
 - we have to apply some kind of performance optimisation measures,
 - we might have to scale our system going forward

2. * Vertical Scaling or Scale-up

- Upgrading our initial single machine to 2x, 3x etc.
- RAM by 2x & SSD by 3x etc.
- Scale up is pocket friendly till a point only
- more you scale up, cost increases exponentially
- Good optimisation as of now

- business is growing, you decided to scale it to 3 more cities and now getting 300 booking per minute

3. * Command Query Responsibility Segregation CQRS

- the scaled up big machine is not able to handle all read/write requests.
- separated read/write operation physical machine wise.
- 2 more machines as replica to the primary machine
- all read queries to replicas
- all write queries to primary
- business is growing, you decided to scale it 2 more cities.
- primary is not able to handle all write requests
- lag between primary and replica is impacting user experience.

6. * Horizontal Scaling or Scale-out

- Sharding - multiple shards
- allocate 50 machines - all having same DB schema - each machine just hold a part of data
- locality of data should be there
- each machine can have their own replicas, may be used in failure recovery.
- Sharding is generally hard to apply. But "No pain, No Gain".
- Scaling the business across continents.

7. * Data Centre Wise Partition

4. * Multi Primary Replication

- why not distribute write request to replica also?
- all machines can work as primary & replica
- multi-primary configuration is a logical circular ring
 - write data to any node
 - read data from any node that replies to the broadcast first

- you scale to 5 more cities and your system is in pain again 2050 request/sec

5. * Partitioning of DATA by Functionality

- what about separating the location table in separate DB schema?
- what about putting that DB in separate machines with primary - replica or multi-primary configuration?
- Different DB can host data categorised by different functionality.
- backend or application layer has to take responsibility to join the results

CAP theorem

- basic and one of the most imp concept in distributed databases
- Useful to know this to design efficient distribution system for your given business logic

Q) Let's first breakdown CAP

- Q) What does the CAP Theorem say?
- a distributed system can only provide two of three properties simultaneously - consistency, availability, partition tolerance. The theorem formalizes the tradeoff between consistency & availability when there is a partition.

Q) CAP Theorem NO SQL Databases

1) CA DB -

- in a consistent system, all node see the data simultaneously. if we perform a read operation on a consistent system, it should return the value of the most recent write operation. The read should cause all nodes to return the same data. all users see the same data at the same time, regardless of the node they connect to. when data is written to a single node, it is then replicated across the other nodes in the system.

Availability :

- when availability present in a distributed system, it means, that the system remains operational all of the time. every request will get response regardless of the individual state of the nodes. This mean that the system will operate even if there multiple nodes down, unlike consistent system. There's no guarantee that the response will be the most recent write operation.

2) CP DB -

- CP DB enables consistency and partition tolerance, but not availability. when a partition occurs, the system has to turn off inconsistent nodes until the partition can be fixed. MongoDB is an example of a CP DB. It's a NoSQL database management system, that uses documents for data storage. It's considered schema-less, which means that it doesn't require a defined database schema. It's commonly used in big data and applications running in different locations. The CP system is structured so that there's only one primary node that receives all of the write requests in a given replica set. Secondary nodes replicate the data in the primary nodes, so if the primary node fails, a secondary node can stand-in. In banking system availability is not as important as consistency, so we can opt for MongoDB DB.

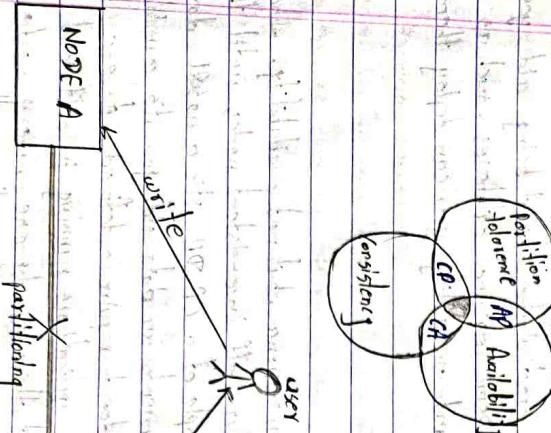
C) Partition Tolerance :

- when a distributed system encounters a partition, it means that there's break in communication between nodes. If a system is partition-tolerant, the system does not fail, regardless of whether message are dropped or delayed between nodes within the system. To have partition tolerance, the system must replicate records across combinations of nodes and networks.

3)

AP DB -

- AP DB enable availability and partition tolerance, but not consistency. In the event of partition, all nodes are available, but they're not replicated for Ex - if a user tries to access data from a bad node, they won't receive the most up-to-date version of the data. When the partition is eventually resolved, most AP databases will sync the nodes to 'ensure consistency' across them. Apache Cassandra is an example of an AP database, it's NOSQL database with no primary node, meaning that all of the nodes remain available.
- Cassandra allows for eventual consistency because users can re-sync their data right after a partition is resolved. For apps like Facebook, we value availability more than consistency, we'd opt for AP databases like Cassandra or Amazon DynamoDB.



4)

ACID & Base Model

- 1) ACID Model
 - collections of operations that form a single logical unit of work are called transactions and the DB system must ensure proper execution of transactions and the ACID DB Transaction model ensures that a performed transaction is always consistent. To explain ACID in more detail and easy way is to underscore through breaking down the acronym, ACID

1) Atomicity -

- this property states transaction must be treated as atomic unit, that is, either all of its operations are executed or none and there must be no state in database where a transaction is left partially completed also the states should be defined either before the execution of the transaction or after the execution of the transaction.

2) Consistency

- the DB must remain in a consistent state after any transaction also no transaction should have any adverse effect on the data residing in the DB and if the DB was in a consistent state before the execution of a transaction then it must remain consistent after the execution of the transaction as well.

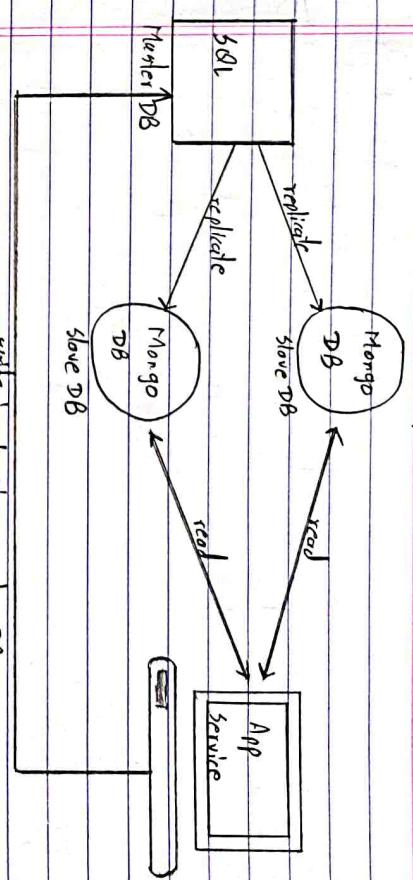
3) Isolation

- in a DB system where more than one transaction is being executed simultaneously and in parallel, the property of isolation states that each one of the transaction is going to be administered and executed as if it is the only transaction in the system also no transaction will affect the existence of any other transaction.

4) Durability

- the DB should be durable enough to hold all its latest updates even if the system fails or restarts so in practical way of saying that if a transaction updates a chunk of data in a DB and commit is performed then the DB will hold the modified data but if the commit is not performed then no data is

The Master-Slave DB concept



Base Model

- the rise in popularity of NoSQL DB's provided a flexible and fluidity with ease to manipulate data and as a result, a new DB model was designed reflecting these properties. The acronym BASE is slightly more confusing than ACID but however, the words behind it suggest ways in which the BASE model is different and acronym BASE stands for

Segregation

- master-slave is a general way to optimise IO in a system where number of requests goes way high that a single DB server is not able to handle it efficiently
- it's a pattern 3 in 12C-19 (DB scaling pattern) Command Query Responsibility Segregation

- 1) Basically Available
 - instead of making it compulsory for immediate consistency, BASE modelled NO-SQL database will ensure the availability of data by spreading and replicating it across the nodes of the DB cluster

Soft State

- due to the lack of immediate consistency, the data values may change over time. the BASE value may change over. the BASE model breaks off with the concept of a DB that obligates its own consistency, delegating that responsibility to developers

Eventual Consistent

- the fact that BASE does not obligates immediate consistency but it does not mean that it never achieves, if however, until it does, the data reads are still possible (even though they might not reflect reality)